Computational Supplement

Frame covariance of CMB lensing in the exact $1+3$ covariant radiation hierarchy

Tim Gebbie

v1.2.0-dev — the source decomposition in two frames, Appendix E made quantitative, and the k grid carried to the damping scale

Last release: v1.1.0. First public analytical reproducibility release: v1.0.0, 5 October 2026
14 October 2026

1 Scientific scope

This supplement records the analytical companion to *Frame covariance of CMB lensing in the exact* $1+3$ *covariant radiation hierarchy*. The bundle is an independent Python reconstruction, from first $1+3$ covariant principles, of the almost-Friedmann–Lemaître results of Gebbie, Dunsby & Ellis, *Annals of Physics* **282**, 321–394 (2000), `doi:10.1006/aphy.2000.6034`. It is archived at `doi:10.25375/uct.34069509`, live since v1.0.0. That is the **concept** DOI and always resolves to the latest version, so it is the one to cite.

It is written bottom-up: not a port of, and not a wrapper on, an existing Boltzmann code, and it inherits no other code's conventions. It reproduces; it does not argue. No new effect is claimed.

Criterion 5b compares against CAMB and CLASS. It does not port them. From v1.2.0 this bundle ships a `reference/` directory holding TT spectra generated by CAMB and CLASS, and that is the only form in which either code appears. **Nothing under `functions/` imports `camb` or `classy`, and nothing ever may**: a reconstruction that imported the codes it is checked against would not be an independent one, and a reader reproducing figure F1 should not need a Fortran toolchain. `scripts/make_reference_spectra.py` is the single caller, it is not a figure stage, and what enters the repository is **data with a stamp** — the spectra, the versions that produced them, every setting, and the derived parameters — never vendored source. Where a published result is carried on this side so that a comparison is like for like, it is taken from the literature and named, not read out of either code. **The external codes are a check on this reconstruction, never an input to it.**

The background expansion is *derived* from the $1+3$ covariant energy constraint rather than transcribed from a printed expansion law, so that the antecedent's own Eqs. (185) and (G.3) are checks on this bundle and not its sources. That ordering is what makes a reconstruction an audit, and it is load bearing: a port would have inherited the misprint recorded below.[1]

Department of Statistical Sciences, University of Cape Town. `tim.gebbie@uct.ac.za`

2 Theory-to-computation map

Table 1 links each published equation the numerics implement to its provenance and to the `.py` file and object implementing it, together with the test that checks it. Table 2 gives the notation and the names

[1] The Gauss constraint at (G.3) prints $+\frac{2}{3}\Theta^2$ where the companion thesis and the Cargèse lectures Eq. (55) carry $-\frac{2}{3}\Theta^2$. In the Friedmann limit the printed sign gives $H^2 = K/a^2 - (\mu+\Lambda)/3$, which for a flat section is $-(\mu+\Lambda)/3$ and so has no solution at positive energy density. A misprint, recorded; the derived background already carries the correct sign, and a test now pins it so that a later edit towards the printed form fails rather than passes.

the quantities carry in the code.

Both are *generated* from `config/implementation-register.toml` by `scripts/make_tables.py`, and `tests/test_register.py` resolves every module, object and test named in them. They cannot drift from the code: a rename breaks the build rather than quietly falsifying a table printed here.

The harmonic normalisation is $Q_{A_\ell} = (-k_{\mathrm{phys}})^{-\ell} \mathrm{D}_{\langle A_\ell \rangle} Q$. Both candidate conventions are run end to end; only this one reproduces the real, opposite-sign form of the external hierarchies.[2]

Nomenclature is kept strictly separate: τ_ℓ are **mode** coefficients and τ_{A_ℓ} are **multipole** coefficients. *Annals II* notes that the distinction is not made in Bardeen-variable treatments. Every released figure plots a multipole mean-square, which holds for general geometries, rather than a mode mean-square, which holds only for almost-Robertson–Walker ones.

[2] *Annals I* prints $(-1)^\ell$ at the corresponding step where its own definition gives i^ℓ. The covariant multipole τ_{A_ℓ} is invariant under the choice, so the misprint never produced a wrong result; the mode coefficients are not invariant, and the two statements are distinct. Both are asserted as tests. A misprint, recorded in `provenance/ANNALS-II-ANTECEDENT-AND-CORRECTIONS-v1.md`.

Table 1: Equations implemented in the numerics, with their provenance and the code that implements them. Equation numbers are those of the *published* papers. Rows with no code are not yet implemented.

Quantity	**Source**	**Equation**	**File**	**Object**	**Checked by**
$\beta_\ell = O^{A_\ell} O_{A_\ell} = (\ell!)^2 2^\ell/(2\ell)!$	Annals I	(24)	`functions/harmonics/weights.py`	`beta`	`test_beta_module_matches_symbolic_exactly`
$\beta_\ell = \frac{\ell}{2\ell-1}\beta_{\ell-1}$, $\beta_{\ell+1} = \frac{\ell+1}{2\ell+1}\beta_\ell$	Annals I	(25)	`functions/harmonics/weights.py`	`beta`	`test_beta_recurrence_descending_general_ell`
$\Delta_\ell = 4\pi\beta_\ell/(2\ell+1)$	Annals I	(119)	`functions/harmonics/weights.py`	`delta_over_four_pi`	—
α_ℓ, defined by $\beta_\ell = \alpha_\ell^{-1}(2\ell+1)$	Annals II	App. F, p. 380	`functions/harmonics/weights.py`	`alpha`	`test_f4_is_direct_substitution_into_f3`
free-streaming weight $w_\ell = (\ell+1)^2/[(2\ell+3)(2\ell+1)]$	Annals II	(F.1), p. 379	`functions/harmonics/weights.py`	`free_streaming_weight`	`test_f1_weight_is_the_mode_recursion_weight_shifted`
covariant recursion $-\dot{\tau}_\ell = (k_{\rm com}/a)[w_\ell\tau_{\ell+1} - \tau_{\ell-1}]$	Annals II	(F.1), p. 379	`functions/harmonics/free_streaming.py`	`free_stream`	`test_the_three_normalisations_are_one_solution`
β-normalised recursion, conformal time	Annals II	(F.3), p. 380	`functions/harmonics/free_streaming.py`	`free_stream`	`test_the_three_normalisations_are_one_solution`
α-normalised recursion	Annals II	(F.4), p. 380	`functions/harmonics/free_streaming.py`	`free_stream`	`test_f4_variable_is_exactly_the_spherical_bessel_function`
free-streaming solution $\alpha_\ell^{-1}\tau_\ell = j_\ell(k_{\rm com}\Delta\eta)$	Annals II	App. F, p. 380	`functions/harmonics/free_streaming.py`	`analytic_projection`	`test_f4_variable_is_exactly_the_spherical_bessel_function`
$\dot{\Theta}_\ell = k[\frac{\ell}{2\ell-1}\Theta_{\ell-1} - \frac{\ell+1}{2\ell+3}\Theta_{\ell+1}]$	Hu & Sugiyama	(6), p. 2601	`functions/harmonics/external.py`	`WEIGHTS`	`test_external_hierarchies_match_the_bundle_at_machine_precision`
$\dot{F}_{\nu\ell} = \frac{k}{2\ell+1}[\ell F_{\nu(\ell-1)} - (\ell+1)F_{\nu(\ell+1)}]$	Ma & Bertschinger	(49), (50)	`functions/harmonics/external.py`	`WEIGHTS`	`test_external_hierarchies_match_the_bundle_at_machine_precision`
$\dot{\Delta}_{T\ell} = \frac{k}{2\ell+1}[\ell\Delta_{T(\ell-1)} - (\ell+1)\Delta_{T(\ell+1)}]$	Seljak & Zaldarriaga	(3d)	`functions/harmonics/external.py`	`WEIGHTS`	`test_external_hierarchies_match_the_bundle_at_machine_precision`

Quantity	Source	Equation	File	Object	Checked by
relative-sign control: the $c^2 = -1$ normalisation, run as a wiring check	Annals II	(F.1), p. 379	`functions/harmonics/free_streaming.py`	`free_stream`	`test_sign_control_fails_in_a_known_shape`
$\int_0^\infty z^{-m} j_\ell^2(z)\, dz$, corrected to carry $[(m/2)!]^2$	Annals II	(199), p. 372	`scripts/derive_ell_range.py`	`bessel_integral_corrected`	`test_corrected_bessel_identity_matches_quadrature`
$\ell_A = \pi \Delta\eta_* / r_s$, bounding criterion 5's range	Annals II	§8.3.2, p. 372	`scripts/derive_ell_range.py`	`acoustic_scale`	—
$(a'/a)^2 = H_0^2[\Omega_r a^{-2} + \Omega_m a^{-1} + \Omega_k + \Omega_\Lambda a^2]$	this work	the $1+3$ energy constraint	`functions/background/flrw.py`	`Background`	`test_the_gauss_constraint_sign_is_the_corrected_one_and_the_printed_one_has_no_solution`
$\mathcal{P}(k) = \mathcal{A}(k\eta_0)^{n-1}$	Annals II	(J.1)	`functions/spectra/transfer.py`	`power_spectrum_j1`	`test_j1_is_constant_at_n_equals_one_and_eta0_drops_out`
$\mathcal{P}(k) = \mathcal{B} k^n T^2(k)$	Annals II	(J.3)	`functions/spectra/transfer.py`	`power_spectrum_j3`	`test_j1_and_j3_differ_by_one_power_of_k`
$T(k) = [1 + (ak + (bk)^{3/2} + (ck)^2)^\nu]^{-1/\nu}$	Annals II	(J.5)	`functions/spectra/transfer.py`	`transfer_function`	`test_b10_the_printed_transfer_function_diverges_where_it_must_vanish`
$\dot\delta_\ell = -n_e \sigma_T \delta_\ell - ikT\{\frac{\ell}{2\ell-1}\delta_{\ell-1} + \frac{\ell+1}{2\ell+3}[1 - \frac{\ell(\ell+2)K}{k^2}]\delta_{\ell+1}\}$	Wilson	(8), p. 4	`functions/harmonics/external.py`	`integrate_wilson`	`test_wilson_reproduces_the_bundle_and_tests_the_phase_convention`
$x_e(z)$, equilibrium Saha	this work	standard; not in the antecedent	`functions/spectra/decoupling.py`	`saha_ionisation`	`test_saha_is_monotone_and_spans_the_full_range`
$\mathcal{V}(k,\eta) \simeq \kappa' e^{-\kappa}$	Annals II	§5.2	`functions/spectra/decoupling.py`	`RecombinationHistory`	`test_the_visibility_is_a_normalised_probability_density`
$\mathcal{D}(\eta_0, k) = \int d\eta\, \mathcal{V} e^{-(k/k_D)^2} \simeq e^{-(k/k_D)^2}$	Annals II	(101), §5.2	`functions/spectra/decoupling.py`	`RecombinationHistory`	`test_the_visibility_is_a_normalised_probability_density`
the integral solution with $e^{-\kappa}$ on gravitational and $\kappa' e^{-\kappa}$ on scattering sources	Annals II	(98), §5	`functions/spectra/decoupling.py`	`Weighting`	`test_the_two_weightings_are_distinct_objects`
the same solution with both weights replaced by $\mathcal{V}$	Annals II	(102), §5	`functions/spectra/decoupling.py`	`Weighting`	`test_the_visibility_weighting_cannot_carry_a_late_isw`

Quantity	Source	Equation	File	Object	Checked by
$\gamma \simeq \frac{k^2}{6\kappa'} \frac{R^2+\frac{4}{5}(1+R)}{(1+R)^2} = k^2/k_D^2$	Annals II	(168), §6	functions/spectra/decoupling.py	diffusion_scale	test_damping_suppresses_small_scales_and_leaves_large_ones
$-C_0 = \frac{5}{2}a\mathcal{B}_2$, $-C_1 = \frac{1}{k}(a\mathcal{B}_1+\kappa' v_B)$, $-C_2 = \frac{1}{k^2}a\mathcal{B}_2$	Annals II	(90), §5	—	—	—
$\tilde{\tau}_1' = -\frac{R'}{1+R}\tilde{\tau}_1 - \frac{k}{1+R}\delta\tilde{T} - k\Phi_A$; $\delta\tilde{T}' = -aH\Phi_A - \Phi_H' + \frac{k}{3}\tilde{\tau}_1$	Annals II	(152), (153), §7.1	functions/spectra/acoustic.py	solve_acoustic	test_the_pair_reduces_to_the_printed_oscillator_exactly
$\delta\tilde{T}'' + \frac{R'}{1+R}\delta\tilde{T}' + k^2 c_s^2 \delta\tilde{T} = -\Phi_H'' - \frac{R'}{1+R}\Phi_H' - \frac{k^2}{3}\Phi_A$	Annals II	(154), §7.1	functions/spectra/acoustic.py	solve_acoustic	test_the_pair_reduces_to_the_printed_oscillator_exactly
$c_s^2 = 1/(3(1+R))$, $R = 3\rho_M/4\rho_R$	Annals II	§6	functions/spectra/acoustic.py	sound_speed_squared	test_sound_speed_limits
$\delta\tilde{T} + \Phi_A \approx [\delta\tilde{T}(0) + (1+R)\Phi_A]\cos(kr_s) - R\Phi_A$	Annals II	(180), §7.1.2	functions/spectra/acoustic.py	sachs_wolfe_at_last_scattering	test_the_integration_converges_to_the_closed_form_180_as_r_is_held_constant
$[\delta\tilde{T} + \Phi_A](\eta_*) \approx \frac{1}{3}\Phi_A(0)(1+3R_*)\cos(kr_s^*) - R_*\Phi_A(0)$	Annals II	(181), §7.1.2	functions/spectra/acoustic.py	sachs_wolfe_at_last_scattering	test_b8_equation_181_sign_is_fixed_by_its_own_limit
$\frac{\beta_\ell \tau_\ell(\eta_0)}{2\ell+1} = S_P j_\ell(k\Delta\eta_*) + \int d\eta\,(S_{\rm DISW} + S_{\rm ISW}) j_\ell(k\Delta\eta)$	Annals II	(176), §7.1.1	functions/spectra/sources.py	temperature_transfer	test_176_with_an_impulsive_source_is_the_free_streaming_projection
$S_P = D(\eta_0, k)\,[(\tilde{\delta}_T + \Phi_A) + \kappa' v_B]$	Annals II	(177), §7.1.1	functions/spectra/sources.py	source_primary	test_176_with_an_impulsive_source_is_the_free_streaming_projection
$S_{\rm DISW} = V e^{-(k/k_D)^2}[\frac{1}{3}k\tilde{\tau}_1 + (\kappa'\tilde{v}_B' + \kappa''\tilde{v}_B)]$	Annals II	(178), §7.1.1	functions/spectra/sources.py	source_doppler	test_the_printed_doppler_bracket_equals_the_contracted_form
$S_{\rm ISW} = V e^{-(k/k_D)^2}[(\Phi_A' - \Phi_H') - aH(\tilde{\delta}_T + 3\Phi_A)]$	Annals II	(179), §7.1.1	functions/spectra/sources.py	source_integrated	test_b6_sign_alternative_is_not_a_no_op
$C_\ell = \frac{2}{\pi}\frac{\beta_\ell^2}{(2\ell+1)^2}\int \frac{dk}{k} k^3 \lvert\tau_\ell\rvert^2$	Annals II	(186), p. 368	functions/spectra/angular.py	cl_mode_route	test_the_harmonic_normalisation_cancels_out_of_186_exactly
$\langle \tau_{A_\ell} \tau^{A_\ell} \rangle = \frac{\beta_\ell}{2\pi^2}\int k^2 dk \lvert\tau_\ell\rvert^2$	Annals II	(187), p. 368	functions/spectra/angular.py	multipole_mean_square	test_188_is_delta_over_2l_plus_1_times_187

Quantity	Source	Equation	File	Object	Checked by
$C_\ell = \Delta_\ell (2\ell+1)^{-1} \langle \tau_{A_\ell} \tau^{A_\ell} \rangle$	Annals II	(188), p. 368	functions/spectra/angular.py	cl_covariant_route	test_the_two_routes_agree_to_machine_precision_on_real_input
$C_\ell^{\mathrm{SW}} = \frac{1}{2\pi} H_0^4 \Omega_0^{1.54} \int \frac{dk}{k^2} P(k) j_\ell^2(k\Delta\eta_*)$	Annals II	(198), p. 371	functions/spectra/sachs_wolfe.py	cl_sachs_wolfe_quadrature	test_the_analytic_reduction_matches_the_quadrature
large-scale closed form, $n = 1$	Annals II	(200), (201), p. 372	functions/spectra/sachs_wolfe.py	cl_large_scale_reduced	test_b5_printed_201_is_short_by_the_comoving_distance
closed-form k-integral kernel at complex m, $\int z^{-m} j_\ell^2(z) dz$	Annals II	(199), p. 372, corrected	functions/spectra/kernels.py	bessel_square_kernel	test_complex_kernel_reduces_to_the_real_identity
primary-term k-integral by power-law decomposition	Annals II	(176) primary term, (198)	functions/spectra/kernels.py	fftlog_primary_integral	test_fftlog_matches_quadrature_on_a_non_power_law_integrand
integrated terms, $\int d\eta\, S(k,\eta)\, j_\ell(k\Delta\eta)$	Annals II	(176), p. 365	functions/spectra/bessel_table.py	line_of_sight	test_line_of_sight_matches_direct_evaluation
truncation condition: four consecutive vanishing harmonics force $\sigma_{ab} = 0$	Annals II	App. E, (E.1), p. 378	—	—	—
C_0, C_1, C_2 from the rank-by-rank match of the integral solution	Annals II	(89)	functions/spectra/coefficients.py	match_coefficients	test_all_three_of_90_are_reproduced_by_the_match
$-C_0 \simeq +\frac{5}{2} a\mathcal{B}_2$, $-C_1 \simeq \frac{1}{k}(a\mathcal{B}_1 + \kappa' v_B)$, $-C_2 \simeq \frac{1}{k^2} a\mathcal{B}_2$	Annals II	(90)	functions/spectra/coefficients.py	c0_from_printed_89	test_b11_is_a_sign_in_89_not_in_90
$-\tilde{C}_1 \approx \kappa' \tilde{v}_B + (\delta\tilde{T} + \Phi_A)$	Annals II	(109)	functions/spectra/coefficients.py	newtonian_coefficients	test_the_newtonian_specialisation_is_equation_109
the slow-decoupling integral solution carried with a generic u^a: two boundary terms and one integrated term	Annals II	(102)	functions/spectra/integral_solution.py	integral_solution	test_the_specialisation_recovers_what_the_pipeline_computes
$(a\tilde{\mathcal{B}}_1)' \approx -a^2 H \tilde{\mathcal{B}}_1 + k(\Phi_A' - \Phi_H') - 2Hak\Phi_A + \frac{1}{3}k^2\tilde{\tau}_1$	Annals II	(106)	functions/spectra/coefficients.py	newtonian_c1_prime	test_t23_splining_the_amplitude_is_wrong_and_the_test_can_tell
$\dot{\sigma}_{\langle ab\rangle} + 2H\sigma_{\langle ab\rangle} - \mathcal{D}_{\langle a} A_{b\rangle} + E_{\langle ab\rangle} \simeq \frac{1}{2}\pi_{\langle ab\rangle}$, and its second derivative	Annals II	(G.12)	functions/spectra/coefficients.py	generic_shear_derivatives	test_the_rank_two_control_passes_symbolically

Quantity	Source	Equation	File	Object	Checked by
C_2, C_2' and C_2'' from $\mathcal{B}_2 = \sigma_{ab}$	Annals II	(90)	`functions/spectra/coefficients.py`	`c2_derivatives`	`test_no_numerical_derivative_survives_in_the_c2_chain`
$\sigma_{ab} = \mathcal{D}_{\langle a} v_{b\rangle}$ in the CDM frame, with $\dot{\sigma}$ and $\ddot{\sigma}$	Annals II	(31) + (G.12)	`functions/spectra/coefficients.py`	`cdm_frame_shear`	`test_B13_the_printed_n_t1_does_not_close_and_the_general_rule_does`
$C_0, C_1, C_2, C_1', C_2''$ in the CDM frame, $A_a = 0$	this work	(89) + (31)-(37) + (106) + (152)	`functions/spectra/coefficients.py`	`cdm_set`	`test_the_frame_shift_of_every_coefficient_is_the_closed_form_one`
$\tilde{\tau}_a = \tau_a - v_a$, as $\tilde{\tau}_1 \to \tilde{\tau}_1 - v/k$	Annals II	(37)	`functions/spectra/frame_transform.py`	`FrameShift`	`test_T27_the_dipole_carries_the_same_inverse_k_as_the_monopole`
$(\rho, p, \mathcal{H}')$ from the Ω's, consistent with $\dot{H} = -H^2 - (\rho + 3p)/6$	Annals II	Appendix G	`functions/spectra/coefficients.py`	`background_fluid`	`test_the_frame_shift_of_every_coefficient_is_the_closed_form_one`

Table 2: Notation, and the names the quantities carry in the code. The mode/multipole distinction is the one that matters most: *Annals II* notes that it is not made in Bardeen-variable treatments.

Symbol	In code	Meaning	Note
Φ_A, Φ_H	`phi_a, phi_h`	the two gauge-invariant potentials entering (177) and (179)	—
$\tilde{\delta}_T$	`delta_T`	covariant temperature perturbation	frame-dependent; $\tilde{\delta}_T + 3\Phi_A$ is the combination in dispute (B-6)
$\tilde{\tau}_1$	`tau_1`	the dipole, which supplies the Doppler term of (178)	—
$\tilde{v}_B$	`v_b`	baryon velocity	—
κ', κ''	`kappa_prime`	differential optical depth and its conformal-time derivative	supplied, never manufactured inside the source terms
V	`visibility`	visibility function, weighting the integrated sources	—
k_D	`damping_k`	diffusion damping scale, entering (178) and (179) as $e^{-(k/k_D)^2}$	—
aH	`conformal_hubble`	a'/a, the background combination (179) carries	derived from the $1+3$ constraint, not from a printed expansion law
S_P, $S_{\rm DISW}$, $S_{\rm ISW}$	`source_primary, source_doppler, source_integrated`	the primary, Doppler and integrated temperature sources of (177)–(179)	individually frame-dependent
$\tau(x, e)$	—	temperature anisotropy as a function of position and direction	—
τ_{A_ℓ}	—	**multipole** coefficients — the covariant objects	valid for general geometries; what every released figure plots
τ_ℓ	`tau, values`	**mode** coefficients	almost-Robertson–Walker only. Conflating these two is how a sign is lost
O^{A_ℓ}	—	$e^{\langle A_\ell \rangle}$, the PSTF basis on the sphere	Annals I (15)
Q_{A_ℓ}	—	harmonic basis tensor, $(-k_{\rm phys})^{-\ell} \mathrm{D}_{\langle A_\ell \rangle} Q$	the adopted normalisation; see criterion 3
β_ℓ	`beta`	$O^{A_\ell} O_{A_\ell}$	Hu & Sugiyama's normalisation
α_ℓ	`alpha`	$(2\ell+1)/\beta_\ell$	Ma & Bertschinger's and Seljak & Zaldarriaga's normalisation
$k_{\rm com}$	`k_com`	comoving wavenumber	**never a bare `k` in this codebase**
$k_{\rm phys}$	`k_phys`	physical wavenumber, $k_{\rm com}/a$	Appendix F's k is comoving; the conventions sheet's is physical
η	`eta`	conformal time, $dt = a\, d\eta$	—
$\Delta\eta_*$	`d_eta`	$\eta_0 - \eta_*$	—
j_ℓ	`spherical_jn`	spherical Bessel function	the free-streaming solution
u^a	—	the frame 4-velocity	carried generic; frames are specialisations made at the end
σ_{ab}	—	shear	σ is the shear and nothing else
A_a	—	acceleration, $\dot{u}_a$	—
Θ_ℓ	—	Hu & Sugiyama's multipole variable	β-normalised

3 The acceptance spine

Appendix F of the antecedent, pp. 379–380, Eqs. (F.1)–(F.4), links the covariant mode recursion to external treatments written in real coefficients. Reproducing that chain tests the harmonic normalisation

without introducing a representation, and it is where this bundle's independence lies.[3]

Five external treatments are read and **four** are drawn — Hu & Sugiyama Eq. (6), Wilson (1983) Eq. (8), Ma & Bertschinger Eqs. (49) and (50), and Seljak & Zaldarriaga Eq. (3d) — each transcribed from its own paper.[4] The fifth is Wilson & Silk (1981) Eq. (7), the earliest of them and likewise read from its own page: at $K = 0$ it *is* Wilson (1983) Eq. (8), so it is credited and not drawn, four panels carrying four independent numerical checks. Hu & Sugiyama and Wilson write the recursion in the β normalisation — the normalisation their Eq. (2) fixes — and the other two in the α normalisation, which is why Appendix F prints both (F.3) and (F.4). A basis phase leaking into the couplings would break at least one of the two matches.

Wilson is the strongest of the four, and not because of his weights, which are Hu & Sugiyama's. He writes the recursion in the *imaginary* convention, with both bracket terms positive under an overall $-ikT$, where the others carry opposite signs under a real k. The two are related by $\Theta_\ell = i^\ell \delta_\ell$, which turns one form into the other exactly. He is integrated complex and as printed, and the phase removed only afterwards; transforming the equation first would assume the identity under test. That the transformed solution is *real* to 10^{-10} is a separate assertion from that it agrees. Wilson therefore tests the phase convention — the subject of §2's footnote and of acceptance criterion 3 — from a source published before any of the others.[5]

4 Conventions, and what the comparison does and does not settle

The same hierarchy is written four different ways in the literature this bundle checks against, and the differences are not cosmetic. Table 3 is the translation, generated from `config/implementation-register.toml` so that it cannot drift from the code, with the test that checks each row.

"RW", "FL" and "FLRW" are three different claims, and this bundle makes the middle one. The names are used interchangeably often enough that it is worth separating them once.

- **Robertson–Walker (RW)** is a statement about *geometry*: the spacetime is spatially homogeneous and isotropic. It says nothing about what drives the expansion, and it follows from symmetry alone — no field equation has been used.
- **Friedmann–Lemaître (FL)** is that geometry *with its dynamics*: the Einstein equations imposed on it, giving the Friedmann and Raychaudhuri equations and hence $a(\eta)$ for a given matter content. FL is RW *plus* the field equations.
- **FLRW** is the two together, and is the usual compound name.

Annals II and this bundle say **almost-Friedmann–Lemaître**, and the choice is not stylistic. The background here is integrated from the $1 + 3$ energy constraint rather than assumed, so the dynamical

[3] The comparison applies the $(2\ell + 1)^{-1}$ division that the text preceding (F.4) specifies and the display omits; (F.4) is otherwise correct as printed. A misprint, recorded.

[4] Ma & Bertschinger print the same recursion twice. Eqs. (49) and (50) are the *massless neutrino* hierarchy, synchronous and conformal Newtonian, and are collisionless; their Eqs. (63) and (64) are the photon hierarchy, carrying the identical recursion plus a Thomson term $-an_e\sigma_{\rm T}F_{\gamma\ell}$. Appendix F's (F.4) is a free-streaming recursion, so (49) and (50) are the comparison and (63) and (64) are not. Published and arXiv numbering agree for this paper, which is not the case for the Annals papers.

[5] Appendix F's citation pairs the name of one paper with the key of another. Both papers are real and both numbers are right for the paper they belong to: Wilson & Silk (1981) Eq. (7) *is* the $\ell > 2$ hierarchy, and Wilson (1983) Eq. (7) is that paper's quadrupole equation while its Eq. (8) is the hierarchy. The prose names Wilson & Silk, the key is Wilson alone and consistent across every citation, and the load-bearing $i^{-\ell}$ remark names Wilson alone — so the match and the phase remark are about Wilson (1983), and the number beside them is correct for the 1981 paper the prose names. One mismatched pairing, not two independent misprints; recorded, and it moves no result.

half is in use from the start: Eqs. (185) and (G.3) are **checks on this bundle rather than inputs to it**. A treatment that only assumed almost-RW would have the geometry and not the expansion history, and could not perform that check.

The distinction is also what the almost-EGS theorem is about, and Sec. 11 depends on reading it the right way round. Almost-isotropy of the measured radiation bounds the *geometry* — an almost-RW conclusion. Getting from there to almost-FL requires the field equations as a second step. Collapsing the two names collapses a two-step argument into an assertion, which is precisely the move that paper is written to avoid.

There are exactly two conventions, and they differ by $c^2 = -1$. Changing the harmonic normalisation $Q_{A_\ell} = \mu^{-\ell}\mathrm{D}_{\langle A_\ell\rangle}Q$ rescales the mode coefficients as $\tau_\ell \to c^\ell\tau_\ell$ with $c = \mu'/\mu$. A recursion $\tau'_\ell = k[A_\ell\tau_{\ell-1} + B_\ell\tau_{\ell+1}]$ then carries $(A_\ell/c,\, B_\ell c)$, so the ratio B/A picks up c^2. At $c = i$ an opposite-sign real bracket becomes a same-sign bracket under an overall i. Hu & Sugiyama, Ma & Bertschinger and Seljak & Zaldarriaga print the first; Wilson prints the second. That is the whole of the difference, and it is asserted symbolically rather than described.

What the match establishes. The ℓ-weights, which are convention-independent, are confirmed by five sources in two normalisations, four of them drawn and numerically independent. The bundle is shown to sit at the same c^2 as Hu & Sugiyama, Ma & Bertschinger and Seljak & Zaldarriaga, and Wilson is shown to sit at the other. *Both* Wilson papers are in the imaginary convention, which settles a standing hypothesis in the negative: no author published this recursion in both conventions, so the split is between authors and not within one.

One source fixes its own phase, and that matters. Ma & Bertschinger define their Legendre expansion with an explicit factor $(-i)^\ell(2\ell+1)$ and state why: it is "chosen to simplify the expansion of a plane wave", so that $\exp(-i\mathbf{k}\cdot\mathbf{x})$ has coefficients $F_\ell = j_\ell(kr)$. Their coefficients therefore sit at $(-i)^\ell$ from the plane wave *by their own definition*, are real, and satisfy the opposite-sign real recursion. For this source the phase is read off the paper rather than assumed, which is exactly the standing the next paragraph says three of the four lack.

What it does not establish, and this release says so. It does **not** establish that $Q_{A_\ell} = (-k_{\mathrm{phys}})^{-\ell}\mathrm{D}_{\langle A_\ell\rangle}Q$ is the right reading of *Annals I*'s definition. Rephasing *both sides* of a comparison by i^ℓ is a change of variable and cannot fail — measured at 6.2×10^{-15} either way — so a match constrains the convention only where the phase applied to the external source is fixed from outside the comparison. For two of the four, no phase is applied, and that is a choice rather than a derivation; Ma & Bertschinger and Wilson are the two that are fixed from outside the comparison.

The exception, and it is the useful one. *Annals II* states, in the line following its flat-case mode functions, that the covariant form "differs by a factor of $i^{-\ell}$ from Wilson since we are using plain mode functions instead of plane waves". That phase is specified by the antecedent, not chosen here. So Wilson's panel in Figure 1 is the one comparison carrying information about the convention, and it is checked in two independent ways: the transformed solution agrees with the reconstruction, *and* it is real to 10^{-10}, which the agreement alone would not show.

Acceptance criterion 3 was therefore **reopened** and restated to the claim the evidence supports. Settling it needs *Annals I*'s definition of Q_{A_ℓ}, not a better match. Nothing computed here depends on it: the covariant τ_{A_ℓ} is invariant under the choice, the mode coefficients are not, and every released figure plots the invariant quantity.

Table 3: How each treatment in the foundational literature writes the same hierarchy, and how this bundle checked it. The two conventions in the literature are related by $\tau_\ell \to c^\ell \tau_\ell$ with $c^2 = -1$ and nothing else. **Where the transformation applied is "none" that is a choice, not a derivation** — which is why acceptance criterion 3 was reopened; see §4.

Treatment	**Variable**	**Bracket weights**	**Phase convention**	**Transformation applied here**	**How it was checked**
Annals II (F.1), Appendix F	τ_ℓ	$\frac{(\ell+1)^2}{(2\ell+3)(2\ell+1)}$ on $\tau_{\ell+1}$, 1 on $\tau_{\ell-1}$	covariant; $Q_{A_\ell} = (-k_{\mathrm{phys}})^{-\ell} \mathrm{D}_{\langle A_\ell \rangle} Q$	—, this is the reference	arithmetic on β_ℓ alone; no representation introduced
Annals II (F.3), Appendix F	$\beta_\ell \tau_\ell$	$\frac{\ell}{2\ell-1}$, $\frac{\ell+1}{2\ell+3}$	real, opposite-sign bracket	none	(F.1)→(F.3) by multiplying through by β_ℓ, as exact rationals
Annals II (F.4), Appendix F	$\alpha_\ell^{-1} \tau_\ell$	$\frac{\ell}{2\ell+1}$, $\frac{\ell+1}{2\ell+1}$	real, opposite-sign bracket	the $(2\ell+1)^{-1}$ the prose states and the display omits	**finding B-2**: printed (F.4) is not yet in the Ma & Bertschinger form
Hu & Sugiyama Eq. (6), p. 2601	Θ_ℓ	$\frac{\ell}{2\ell-1}$, $\frac{\ell+1}{2\ell+3}$	real, opposite-sign bracket	none — *a choice*, see Q1	integrated as printed from its own paper, never through Appendix F
Wilson Eq. (8), p. 4	δ_ℓ	$\frac{\ell}{2\ell-1}$, $\frac{\ell+1}{2\ell+3}$	**imaginary**; both bracket terms + under an overall $-ikT$	$\Theta_\ell = i^\ell \delta_\ell$ — **fixed by the antecedent**, which states the $i^{-\ell}$ itself	integrated *complex and as printed*; the phase removed only afterwards, and the result asserted real to 10^{-10} separately from that it agrees
Wilson & Silk Eq. (7), p. 15	δ_ℓ	$\frac{\ell}{2\ell-1}$, $\frac{\ell+1}{2\ell+3}$	**imaginary**, as Wilson 1983 — *not* the real form	none required: at $K = 0$ it is Wilson 1983 (8) exactly	read from its own page. The **earliest** of the five and a fifth source for the weights, but **not** a fifth numerical check — the same equation cannot check itself
Ma & Bertschinger Eqs. (49), (50)	$F_{\nu\ell}$	$\frac{\ell}{2\ell+1}$, $\frac{\ell+1}{2\ell+1}$	real, opposite-sign bracket	none — *a choice*, see Q1	integrated as printed from its own paper; synchronous and conformal Newtonian gauges

Treatment	**Variable**	**Bracket weights**	**Phase convention**	**Transformation applied here**	**How it was checked**
Seljak & Zaldarriaga Eq. (3d)	$\Delta_{T\ell}$	$\frac{\ell}{2\ell+1}$, $\frac{\ell+1}{2\ell+1}$	real, opposite-sign bracket	none — *a choice*, see Q1	integrated as printed; also the method CMBFAST implements, cited for criterion 5b

The second table, and what its two thesis columns are for. Table 3 translates between the treatments this bundle is *checked against*. Table 4 does something different: it writes each coupling of this reconstruction in the conventions of Maartens, Gebbie & Ellis (1999), Challinor (2000), the thesis in **both** of the conventions it uses, Beneke & Fidler (2010) and Pitrou (2009). Every cell is derived from a named line of the normative sheet, and no cell is blank: where a treatment has no established counterpart the cell says so.

Two things in it are worth reading rather than skimming. **The thesis takes two columns and they differ on exactly the kinematic rows** — its $\dot{\mathcal{T}}$ form moves every kinematic term to the right-hand side, flipping the overall sign against its K/Π form. Elsewhere the two agree and the table says so in as many words, because a difference in wording where there is none in substance is precisely how two conventions get conflated while appearing distinct. And **the Beneke & Fidler and Pitrou columns are mostly "not established"**: the sheet records both as unsettled at the level of hierarchy coefficients, Beneke & Fidler working in (ℓ, m) tetrad form and Pitrou in Poisson gauge, and neither has had a hierarchy coefficient reproduced here. That is the honest content of those columns and it is printed rather than left to inference.

Table 4: The couplings of this reconstruction written in each external convention. Every cell is derived from the line of `provenance/conventions.md` named in the *Sheet* column; no cell is blank, and a treatment with no established counterpart says so rather than being left to inference. **The thesis takes two columns and they differ on exactly the kinematic rows** — the $\dot{\mathcal{T}}$ form moves every kinematic term to the right-hand side, flipping its overall sign against the K/Π form. On the remaining rows the two agree, and the table says *as Th-K* rather than restating them: a difference in wording where there is none in substance is how two conventions get conflated while appearing distinct.

Coupling	**Sheet**	**MGE99**	**CL**	**Th, K/Π**	**Th, $\dot{\mathcal{T}}$**	**BF**	**P09**
Rank raising	`C4.2`	identity	identity in form; D_a unchanged, the sign sits in k	identity	as Th-K	not established: BF work in (ℓ, m) tetrad form, no coefficient reproduced	not established: Pitrou works in Poisson gauge, no hierarchy coefficient checked
Rank lowering	`C4.4`	identity	identity in form; CL (2.25) prints $+\mathrm{D}^b I_{bA_\ell}$, reproduced from MGE99 (71)	identity; Th App. D footnote to (D.44) at $\ell = 1$ agrees	as Th-K	not established	not established
Free-streaming mode recursion	`C4.5`	identity	identity in form; CL (2.25) free streaming reproduced from MGE99 (71)	identity	as Th-K	not established	not established
PSTF contraction on the sphere	`C6.1`	identity (one source: EMT 1983, via C6.4)	identity; CL expands in the same O^{A_ℓ}	identity, Th (1.20)	as Th-K	not established: BF use ${}_sY_{\ell m}$, not a PSTF basis	identity at the basis level: Pitrou (1.28) expands in n^{A_ℓ}
PSTF self-contraction β_ℓ	`C5.2`	identity	identity; CL absorbs it into Δ^C_ℓ, see the row below	identity, Th (1.23)	as Th-K	no counterpart: (ℓ, m) normalisation carries $\sqrt{4\pi/(2\ell+1)}$ instead	identity, Pitrou (1.29) in the equivalent $4\pi\ell!/(2\ell+1)!!$ form
Angular measure Δ_ℓ	`C5.5`	identity, MGE99 (47), same value	$\Delta^C_\ell = 4\pi(-2)^\ell(\ell!)^2/(2\ell+1)! = (-1)^\ell\Delta_\ell$, CL (2.23)	identity	as Th-K	no counterpart in this form	identity, Pitrou (1.29), verified symbolically

Coupling	**Sheet**	**MGE99**	**CL**	**Th, K/Π**	**Th, $\dot{\mathcal{T}}$**	**BF**	**P09**
Shear coupling at $\ell+2$	`C9, T-MGE and T-Th-T`	$-(\ell+2)$ as printed in (71), (88), (90), (99) — **defect E9**; $+(\ell+2)$ is correct, and MGE99 (89) at $\ell=2$ prints $+\frac{8}{15}\rho\sigma$, agreeing with the correction	$-\frac{8}{15}\delta_\ell^2 I\sigma$ gives $+(\ell+2)$, independently of MGE99	$-(\ell+2)$ in Th (2.73) — carries E9	$+(\ell+2)$ **on the RHS** in Th (2.84): every kinematic term flips overall sign against (2.73) because it is moved across. **This row is where the two thesis columns differ**	not established: BF put the observer at rest, so they carry no shear coupling	not established
Acceleration coupling at $\ell\pm1$	`C9, T-CL and T-Th-T`	identity	$A_a \to -A_a$ under $g \to -g$ at fixed u^a (CL footnote to §2), the signature difference and nothing more	identity	overall sign flipped with the rest of the kinematic terms — **differs from Th-K**	not established: observer at rest	not established

5 The generic source derivatives

*Decision **S26**. Written from `provenance/GENERIC-SOURCE-DERIVATIVES-v1.md`, which carries the working and the open items.*

What (102) needs, and what the paper never wrote. (102)'s integrand is $C_0 - C_1' + C_2''$ and its boundary terms carry $C_2'(\eta_*)$. With (90),

$$-C_0 \simeq +\tfrac{5}{2}a\mathcal{B}_2, \qquad -C_1 \simeq \tfrac{1}{k}(a\mathcal{B}_1 + \kappa' v_B), \qquad -C_2 \simeq \tfrac{1}{k^2}a\mathcal{B}_2,$$

so the solution needs $(a\mathcal{B}_1)'$, $(a\mathcal{B}_2)'$ and $(a\mathcal{B}_2)''$ **for arbitrary** u^a. *Annals II* prints these only after the Newtonian specialisation has removed the shear, which is exactly the specialisation this bundle declines to make early. The derivations are therefore assembled here, and this section records what is closed, what is assembled, and what is open.

The three sources, by rank, with no frame chosen. From §4's integrated Boltzmann equation $\dot\tau + e^a \mathrm{D}_a\tau + \mathcal{B} \simeq C[x,e]$,

$$\mathcal{B} = -\tfrac{1}{3}\mathrm{D}^a\tau_a + (\mathrm{D}_a \ln T + A_a)e^a + \sigma_{ab}e^a e^b,$$

giving $\mathcal{B}_0 = -\frac{1}{3}\mathrm{D}^a\tau_a$, $\mathcal{B}_1 = \mathrm{D}_a \ln T + A_a$ and $\mathcal{B}_2 = \sigma_{ab}$. **All three vanish in the background**, so all three are gauge-invariant at first order by Stewart–Walker — which is the paper's own *"note that these are CGI with respect to u^a"*.

The scale-factor rule, and why the footnote to (105) carries $2H$. A rank-n PSTF source's natural variable carries a^n: the paper writes $(a\tilde{\mathcal{B}}_1)'$ at (106) and $(a^2\sigma_{ab})^{\cdot}$ in Appendix G. So

$$\dot{\mathcal{B}}_1 + 2H\mathcal{B}_1 = \frac{1}{a^2}(a^2\mathcal{B}_1)^{\cdot},$$

and the rank-two object built from a rank-one source, $\mathrm{D}_{\langle a}\dot{\mathcal{B}}_{b\rangle}$, carries the rank-two power. **It is not a commutator correction bolted on; it is the rank showing up in the measure.** That is conventions row **C3.1a** read in the direction nobody had needed, and it is checkable: under C3.1, C3.2 and C4.2 the mode functions are time-independent at every rank, so *all* the H's come from the a^n and none from the Q's — which **predicts (106)'s $2H$ before (106) is consulted**. Implementing it the other way, by splining a stored amplitude, is wrong by a median of 100%: finding **T-23**, and the reason `integral_solution` takes every derivative as an input rather than differentiating one.

Rank 2: closed. $\dot{\mathcal{B}}_2$ *is* in the paper generically, in Appendix G, and closes to second order **with an exact symbolic control** that fails when the rank-two term is removed — so the check is one that can fail, which S27's C rung requires and finding **T-24** exists because an earlier one could not.

Rank 1: assembled, and carrying one term more than (105) prints. $\dot{\mathcal{B}}_1$ is assembled from printed equations. Three of its four terms close exactly against (105)/(106) — $\mathcal{B}_0$, $\dot\Phi_A$ and $-2H\Phi_A$, the last only because both of its contributions are present — and substituting the expansion–potential relation of §*newtg* reproduces every term (105) prints. The generic form then carries **one extra**:

$$(\star) \qquad -H\,\mathrm{D}_{\langle a}\tilde{\mathrm{D}}_{b\rangle}\ln T.$$

This is the temperature-gradient half of $-H\tilde{\mathcal{B}}_a$, whose acceleration half *does* survive into (105). **Two readings, and this bundle chooses neither.** It may be the small-scale approximation — §6 drops all terms of $O(\epsilon\epsilon_H)$ and $(\star)$ has that shape — except that so does $-2H\mathrm{D}_{\langle a}\mathrm{D}_{b\rangle}\Phi_A$, which is kept. Or it

may be an omission, in which case (105), (106) and what descends from them are short a term of the same order as ones they keep. It is first order and it does not obviously vanish, which is enough to make the question real; settling it needs §2.5's $O(\epsilon\epsilon_H)$ bookkeeping read against each term of (105), not a judgement about which is negligible. **Logged as open, not filed as a finding against the paper.**

Rank 0: not needed. $\dot{\mathcal{B}}_0$ is not required by (102) at all — C_0 carries $\mathcal{B}_2$, not $\mathcal{B}_0$. An earlier plan said otherwise and was wrong.

Status. The derivation is complete for what the integral solution needs. What remains open is ($\star$) and the dropped $a\mathcal{B}_0\delta_{\ell 0}$ row of (89) — **both questions about the antecedent rather than gaps in this assembly**, and both recorded rather than resolved by assumption.

6 Frames and specialisation

"Frame" is being asked to do two jobs, and they are different objects. The word is used in this literature for something specified at a single event and for something specified throughout a region, and the two are not interchangeable. The PI's enumeration, recorded here because the rest of this section depends on it:

Object	What is specified	Where
Instantaneous observer rest frame	A unit timelike vector $u^a(p)$, determining a local time direction and an orthogonal rest space	At an event p
Local orthonormal frame	A tetrad $e_{\hat{\alpha}}{}^a(p)$, with $e_{\hat{0}}{}^a = u^a$, also choosing three spatial axes	At an event
Observer worldline	A timelike curve with tangent u^a	Along one observer's history
Threading	A smooth timelike congruence, represented by a unit tangent field $u^a(x)$	Throughout a spacetime region
Adapted orthonormal frame field	A tetrad field whose timelike leg follows the congruence	Throughout that region

Nothing in this bundle is a tetrad. No calculation here chooses spatial axes: every quantity is a PSTF tensor built from u^a and $h_{ab} = g_{ab} + u_a u_b$, so rows two and five are named only to be set aside. What the bundle actually uses is rows one and four, and it uses both:

- the **receiving observer's frame** — a u^a at one event, reception at η_0 — which is what the *measurement* is referred to, and in particular what the dipole is a statement about;
- the **CDM threading** and the **Newtonian threading** — unit tangent fields $u^a(x)$ on the whole region between last scattering and reception — which is what the *propagation* integral is carried along, and what σ_{ab}, A_a and q_a are properties of.

The general transformations (31)–(37) relate two *threadings*; the specialisations (40)–(45) stipulate one. $\tilde{\sigma}_{ab} = 0$ and $\tilde{A}_a = 0$ are conditions on a congruence throughout a region and are not statements about any single observer, which is why neither can be imposed "at the observer" and why the Newtonian one is admissible only in restricted cases. Figure F7's two panels are exactly this pair kept apart: the $\ell \geq 2$ panel compares two *threadings* carrying the same observable, and the dipole panel is where the *receiving observer's frame* enters, because the dipole is the one multipole that is a statement about u^a at reception rather than about the congruence that brought the radiation there.

The words this bundle uses, and why none of them is retired. *Threading* names the congruence **as an object**. *Receiving observer's frame* names the local u^a **at reception**. *Frame* continues to name a **choice of observers** and, following *Annals II* §4.1's own usage, each named choice — the energy frame, the zero-acceleration (CDM) frame, the Newtonian frame. Every one of those is a threading in the first sense: the words sit at different levels and are not competing. Decision **S29a**.

No identifier is renamed. The code says `"cdm"` and `"newtonian"`, and *Annals II* prints "frame" for exactly these objects, so the code agrees with the source it reconstructs. The rule that a word naming two things should have one of them renamed is satisfied here by **naming the two objects**, which is what the table above does. And the rename operation is itself the hazard: a blanket substitution in this bundle inverted the meaning of diagnostic D7, whose subject is the CDM *threading* measured against the **total-energy** frame — the one place both senses appear in a single sentence, and therefore precisely what a substitution cannot see. Finding **T-33**. **A rename a regular expression can perform is a rename whose exceptions nothing will catch.**

Specialisation. The solution is carried with a **generic** u^a. The explicit frames appear as choices made at the end, not assumed at the start: the Newtonian frame $\tilde{\sigma}_{ab} = 0$, which is the specialisation *Annals II* itself makes at p. 365 in order to remove the shear terms, and the **CDM frame** $\tilde{A}_a = 0$. Carrying the generic frame and specialising late is the difference the $1 + 3$ covariant approach buys; a gauge-fixed treatment fixes the frame before the hierarchy is solved and cannot exhibit it.

"CDM" names two different things, and this supplement keeps them apart. *CDM the model* is a cosmology — flat, matter dominated, no Λ — and it is one of the two models this release delivers a spectrum for. *The CDM frame* is a choice of frame u^a: the congruence in which the cold dark matter is at rest. The two are independent. The ΛCDM *model* is computed in the same *frame*, and a change of model is not a change of frame. In the code the models are named `CDM_MODEL` and `LCDM_MODEL` and the frames `"cdm"` and `"newtonian"`, so the two cannot be confused at a call site.

The CDM frame is *not* the total-energy frame, and an earlier version of this supplement said it was. The definition used here is exact in general relativity and needs no approximation: cold dark matter is **pressureless and collisionless**, so it is acted on by no non-gravitational force and its worldlines are **geodesics**. The CDM frame is therefore the frame with $\tilde{A}_a = 0$, which is what *Annals II* §4.1 calls the *zero acceleration frame*, listing it as the second of three useful choices — the first being the **total-energy frame** $q_a = 0$ and the third the Newtonian frame. They are **three stipulations, not two**, and the paper gives the reason this calculation wants the second: it is "geodesic right through tight-coupling, slow decoupling and into the free-streaming era", whereas the energy frame is recommended there for *two* coupled species and this is a three-species problem that crosses decoupling.

The two agree only in the limit of matter domination. Setting $q_a = 0$ in a mixture gives $\rho_{_C} v^a_{_C} + \frac{4}{3}\rho_{_R} v^a_{_R} + \ldots = 0$, so in the CDM rest frame the total flux is *not* zero — the baryons and the photons are moving — and the energy frame tilts away by $(\rho_b + \frac{4}{3}\rho_{_R})/\rho_{_C}$. Measured in diagnostic D7, that is **0.24** at last scattering for this release's background and **0.49** for the stamped Planck column, with the baryons carrying 37% of it. It is largest exactly where the acoustic physics is, and falls afterwards on both counts, since $\rho_{_R}/\rho_{_C}$ decreases and the baryons stop moving with the photons once they decouple. Recorded as finding **T-29**; the amendment is decision **S8a**.

The one-line reconciliation, and it is the PI's. The CDM frame is the energy frame *of the CDM component*. For pressureless dust $T^{C}_{ab} = \rho_{_C} u^{C}_a u^{C}_b$, so its flux relative to a frame $u^a = u^a_{_C} + v^a$ is $q^{C}_a = \rho_{_C} v_a$ to first order, and $q^{C}_a = 0$ gives $v_a = 0$ — the CDM at rest, exactly. So the two conditions are the same statement about one species, and the conflation this supplement previously carried was the

dropped qualifier: "energy frame" unqualified means the *total*, and the total includes the baryons and the photons, which are moving. Saying *of what* settles it.

The frames, and what each one fixes. *Annals II* §4.1 enumerates four (its sentence says "three"); the baryon and radiation rows below are frames the paper uses elsewhere and are included because they are the ones a reader is most likely to reach for next.

Threading	Defining condition	What it establishes
CDM *(used here)*	$u^a = u^a_{\scriptscriptstyle C}$, equivalently $\tilde{A}_a = 0$	CDM is at rest; its congruence is geodesic in the conserved-dust model, which is exact rather than approximate. *Annals II* calls it the *zero acceleration frame*.
Total energy	$q_a^{\rm total} = 0$	Total momentum density vanishes; **individual species can still move**. Agrees with the CDM frame only to $(\rho_b + \frac{4}{3}\rho_{\scriptscriptstyle R})/\rho_{\scriptscriptstyle C}$.
Baryon	$u^a = u^a_{\scriptscriptstyle B}$	Baryons are at rest; pressure and scattering can accelerate them, so this frame is **not** geodesic.
Radiation energy	$q^\gamma_a = 0$	The radiation brightness dipole vanishes; **higher multipoles can remain**.
Newtonian *(used here)*	$\sigma_{ab} = 0 = \omega_a$, where admissible	Useful for the scalar almost-FLRW comparison and for the Bardeen link; generally accelerated. Admissible only in restricted cases.
Constant expansion	$\mathrm{D}_a\Theta = 0$, where admissible	Removes the spatial expansion gradient relative to that frame; sometimes useful on small scales. Not used here.

Reading down the second column is the point: **each row fixes a different thing, and no two of them are the same condition**. A frame is chosen and then named, not inferred from what it is called.

One of these frames contains observers and the other does not. This is the reason the distinction is worth the space, and it is where the $1+3$ covariant account differs from a metric one in substance rather than in notation.

The **CDM frame is a congruence of observers**. It is defined by a property of the matter — pressureless, collisionless, therefore geodesic — and the definition is exact in general relativity, not a first-order statement. We observe from inside it: a detector in a bound structure is comoving with the matter to the accuracy linear theory works at, and after decoupling the baryons have fallen into the same flow as the cold dark matter, so the baryon and CDM frames coincide at η_0 where the measurement is made. *It is "us".*

The **Newtonian frame is not**. $\tilde{\sigma}_{ab} = 0$ is a condition imposed to simplify, not a property any matter has, and *Annals II* §4.1 says so in terms: the frame "is only consistent in restricted cases" and is "particularly useful in making comparisons with much of the analytic literature". It is the covariant counterpart of what conformal Newtonian gauge does in a metric treatment — admissible about almost-FLRW, and a device.

Why this costs nothing at first order and will not stay free. τ_{A_ℓ} for $\ell \geq 2$ is frame-invariant at first order, so the *observable* does not care which frame carried it there, and a calculation performed in the Newtonian frame returns what an observer in the CDM frame measures. That is the whole licence for the Bardeen-variable route, and it is a first-order licence.

At $O(\epsilon^2\ell)$ it lapses. The couplings v1.5.0 adds are products of two first-order quantities, and a product of two frame-dependent quantities is not frame-invariant just because each factor's leading frame-dependence cancels. Then the frame has to be the one the observer is actually in, and that is the CDM frame, not a device.

This is what figure F5 is for. Its third panel is the amount of machinery that is frame-dependent — in the Newtonian frame one rank carries everything, in the CDM frame all three are occupied — and its first two panels are the first-order invariance that currently makes that difference harmless. The figure shows both at once, which is why it is the right figure to carry into v1.5.0: it is a picture of exactly what stops cancelling.

What the external codes use, and why criterion 5b has to check it. Criterion 5b compares this bundle against CAMB and CLASS. That comparison is only meaningful if the frame each code evolves in is known, because a quantity compared across two frames is not the same quantity. The correspondence below is **verified where it could be, and recorded at its real strength where it could not**; verifying it is a **precondition of 5b**, not an afterthought.

CLASS is verified from its own source, at `classy v3.4.1`. `perturbations.c` carries the comment *"cdm velocity vanishes in the synchronous gauge"*, and `index_pt_theta_cdm` is allocated as an evolved variable only when `gauge == newtonian`: in synchronous gauge θ_C is not evolved because it is zero. The decisive line is CLASS's own error message — *"In the synchronous gauge, it is not self-consistent to assume no CDM: the latter is used to define the initial timelike hypersurface"* — which says, in the code's own words, that **its synchronous gauge is realised by the CDM congruence**. That is the CDM threading of the table above, $u^a = u^a_C$. Both `gauge = synchronous` and `gauge = newtonian` were run and computed.

CAMB is corroborated but not verified to the same standard. Its Python layer exposes `v_newtonian_cdm` and `v_newtonian_baryon` as *derived outputs*, notes that *"L=1 dipoles are Newtonian Gauge"*, and carries an explicit `frame` argument for sources that are not frame-invariant — which is what a code evolving in one frame and converting on output looks like. That is corroboration from an interface, not a reading of the evolution equations, and the decisive statement lives in CAMB's own covariant formulation paper rather than in anything this bundle inspected. It is recorded at that strength and no higher.

Code	**Gauge used for evolution**	**Corresponding $1+3$ frame**
CLASS *(default)*	Synchronous, with CDM velocity divergence $\theta_C = 0$	CDM-comoving: $u^a = u^a_C$, $A_a = 0$; shear generally nonzero.
CLASS *(optional)*	Conformal Newtonian / longitudinal, selected with `gauge = newtonian`	Newtonian: scalar shear and vorticity vanish; acceleration generally nonzero.
CAMB	Covariant equations evaluated in the CDM frame, equivalent to CDM-comoving synchronous gauge	CDM-comoving: $v_C = 0$, $A_a = 0$; shear generally nonzero.
Bolt.jl	Conformal Newtonian gauge, implemented as `BasicNewtonian`	Newtonian: CDM carries an explicitly evolved peculiar velocity relative to this frame.

Two things follow if the table holds. First, **CAMB and CLASS's default both evolve in the frame this bundle uses**, which is the convenient case: 5b compares like with like without a frame transformation in between. Second, **CLASS's `newtonian` option and Bolt.jl evolve in the other one**, so a comparison against either is a comparison across frames and must either transform or restrict itself to the frame-invariant $\ell \geq 2$ multipoles. Figure F5 is the measurement of what that costs: the totals agree and the decompositions do not.

That CAMB is described as using covariant equations in the CDM frame is the part most worth checking rather than assuming, because if true it makes the comparison closer than a gauge-versus-frame translation would suggest, and a comparison that is closer than it should be is the kind this bundle has already been wrong about once.

For readers coming from the metric approach: a frame is not a gauge. This is the distinction most easily lost in translation, and it is worth stating plainly because the two freedoms are often fixed together in a metric-based treatment and so look like one freedom.

- A **gauge** is a choice of coordinates — equivalently a choice of correspondence between the perturbed spacetime and the background. In the $1+3$ covariant approach it is **already dealt with**: every variable used here vanishes in the background, so by the Stewart–Walker lemma it is gauge-invariant by construction. There is no residual gauge mode to fix and none to worry about.
- A **frame** is a choice of the timelike congruence u^a along which the fluid and radiation variables are defined. It is a *physical* choice — which observers — and it survives after the gauge question has been answered. It is the only freedom left, and it is the one this release exhibits.

So the familiar correspondences are correspondences and not identities. Synchronous gauge is usually *completed* by frame along the CDM, and conformal Newtonian gauge is usually *completed* by frame along a shear-free congruence; in each case the gauge condition fixes coordinates and a separate stipulation fixes the observers. A statement that is frame-dependent here will appear in a metric treatment as a statement that depends on which gauge was chosen, and the two dependences are not the same dependence. That is why figure F5 can show the *total* common to two frames while the *decomposition* differs: the split into ranks is frame-dependent and the observable is not, and neither fact is about coordinates.

This is also what fixes the status of the next release. Because v1.0.0 is the general case, **v1.5.0's restricted case is a recovery and a check** — it must return v1.0.0's numbers, and that requirement is the regression test rather than a new result.

The covariance claim is established symbolically, and that is the stronger form. Acceptance criterion 2 asks that a change of frame $u^a \to \tilde{u}^a = u^a + v^a$ return the rank-one bracket unchanged up to $\dot{v}_a$. It does, *identically* in exact arithmetic rather than to a tolerance, and the cancellation is exhibited rather than quoted:

- $\tilde{\mathrm{D}}_{\underline{a}} \ln T = \mathrm{D}_a \ln T - H v_a$, the extra term arising because the projector moves with the frame and $\overline{\ln T} = -H$ in the background;
- $\tilde{A}_a = A_a + \dot{v}_a + v^b \nabla_b u_a$, where $v^b \nabla_b u_a$ is formed *in full* — Hv_a together with $\sigma_{ab} v^b$ and $\epsilon_{abc} v^b \omega^c$ — and the latter two leave at $O(\varepsilon^2)$ by series expansion rather than by hand;
- so $\tilde{\mathcal{B}}_1 = \mathcal{B}_1 + \dot{v}_a$ with **no shear and no** Hv. It is the *same* H twice, which is why the cancellation is structural and not a coincidence of two fitted coefficients.

The curvature half. The deflection depends only on $\Phi_A - \Phi_H$, whose trace-free screen-projected Hessian is the electric Weyl tensor, $E_{ab} \simeq \frac{1}{2}\mathrm{D}_{\langle a}\mathrm{D}_{b\rangle}(\Phi_A - \Phi_H)$. Since E_{ab} vanishes in the background, the Stewart–Walker lemma makes it frame-invariant at first order, and therefore so is $\Phi_A - \Phi_H$. Both potentials shift by the same $(v' + \mathcal{H}v)/k_{\mathrm{com}}$, and the residual $\tilde{E}_{ab} - E_{ab}$ is identically zero — **with a negative control**: a rule leaving Φ_H fixed, or flipping its sign, leaves a residual proportional to $(v' + \mathcal{H}v)$ and fails. The invariance is therefore forced rather than chosen to make a residual vanish. Only the dipole moves, $\tilde{\tau}_a \simeq \tau_a - v_a$; every τ_{A_ℓ} with $\ell \geq 2$ is unchanged, which is the standard any frame dependence surviving at $\ell \geq 2$ must meet: it is an artefact of the reduction, not physics.

The energy frame is derived, not stipulated. Imposing $\tilde{A}_a = 0$ gives $v' + \mathcal{H}v = -k_{\mathrm{com}}\Phi_A$ — the Euler equation of a pressureless geodesic fluid, which is what the cold dark matter obeys. The normalisation is corroborated independently: in the baryon-dominated limit $R \to \infty$, Eq. (152) reduces to that same equation, so the photon dipole and the frame velocity carry one normalisation.

What this release does *not* do, stated because the earlier draft of this section claimed it. It does not compute an assembled spectrum in two frames and draw the difference. Eqs. (177)–(179) are **Newtonian-frame expressions**: they descend from (106), which the corrections record flags in those words, and the Newtonian frame carries its whole $O(\ell)$ source in the acceleration pair while the CDM frame carries its whole source in the shear pair. Transforming the four input fields and re-evaluating the same equations therefore deletes the coupling that was carrying the source without supplying the one that replaces it: measured, the transfer function moves by a factor of about five, and under the opacity weighting the source collapses to two parts in a thousand. That is not a frame dependence of the sky. The figures which would draw it, and the convergence figure which would test footnote 29 on p. 365 in both frames, require the generic-u^a hierarchy and are deferred with it.

One coefficient is changed for criterion 5b, and only there. The photon quadrupole's Thomson collision term carries $\Pi = \Theta_2 + E_0 + E_2$, so polarisation feeds back into the temperature hierarchy and strengthens the diffusion damping. CAMB and CLASS both include this **structurally**: it is not an output option, and neither can be asked to leave it out — CLASS fails outright in `perturbations_solve` below `l_max_pol_g = 2`, because Π needs E_2, and CAMB exposes no switch for it. Truncating CLASS's polarisation hierarchy from 10 to 4 moves the acoustic band by 0.043%, which bounds the hierarchy's *depth* and not the *feedback*.

The only way to compare like for like is therefore to carry the polarised coefficient here, for that comparison alone. In (dif-damp)'s quadrupole term *Annals II* prints 4/5; Hu & Sugiyama's unpolarised form is 8/9; with polarisation it is 16/15. Measured on `PLANCK_MODEL`:

	coefficient	$k_D(\eta_*)$	ratio
Annals II (dif-damp), as printed	4/5	**619.9**	1.000
Hu & Sugiyama, unpolarised	8/9	588.1	0.949
with polarisation	16/15	**536.9**	**0.866**

a 13.4% reduction in k_D, which takes $e^{-(k/k_D)^2}$ from 0.660 to 0.574 at $k = 400$ and from 0.189 to 0.109 at $k = 800$. It is one of the dominant terms in criterion 5b's residual.

4/5 **remains the default and the released value.** Every figure, diagnostic and released number in this bundle uses it, because the bundle reproduces and 4/5 is what the page prints. 16/15 is reached only through an explicit argument, by criterion 5b. Decision **S31**, finding **T-39**. *Annals II*'s 4/5 matches neither standard coefficient, which is recorded as a **question** — it may be a convention in how R and the factor of 6 are distributed — and not as a correction.

7 Comparison and its range

The C_ℓ comparison is Sachs–Wolfe only: (196) descends from (181), in which the acoustic modulation $\cos(kr_s)$ is dropped by taking $r_s^* \to 0$, and §8.3.2 sets the transfer function to unity. Its accuracy is therefore bounded by the term that was dropped. The range is derived, not chosen: $2 \le \ell \le 20$ at 5%, generated by `scripts/derive_ell_range.py`.

The wider comparison cites CMBFAST for its *method* — line-of-sight integration, Seljak & Zaldarriaga, the same reference the antecedent's own bibliography carries — and compares against **CAMB** and **CLASS**, because the Fortran 77 original is unmaintained and effectively unrunnable. Saying so is more honest than claiming a comparison against a code nobody can run. It is reported in three bands, whose failure modes differ — $2 \le \ell < 30$, cosmic-variance dominated and meaningful only binned; $30 \le \ell < 1000$, the acoustic peaks, where a hierarchy or recursion error actually shows; and $1000 \le \ell \le 2000$, the damping

tail, where C_ℓ is small and fractional differences blow up on a vanishing signal. The comparison is made against the corrected closed form.[6][7]

7.1 Criterion 5b's tolerance, and what it is set against

The tolerance is *measured* rather than chosen; what it is measured against changed once it was run. The original rule took the CAMB–CLASS mutual difference as the floor, on the ground that no agreement can be claimed tighter than two trusted independent codes agree with each other. That floor is 0.370% across the acoustic peaks, and this reconstruction stands at 137% there. It was never going to meet it: (102) is a slow-decoupling solution with no polarisation hierarchy, no neutrino anisotropic stress and a visibility-weighted line of sight, set against two full Boltzmann codes. **A criterion that can only be failed is not a criterion.**

What *is* gateable is the thing the antecedent writes twice and recovers symbolically but never evaluates. Its integral solution appears as (102), carried with a generic u^a and specialised afterwards, and again as the source assembly (177)–(179), written in the Newtonian frame from the start. The paper states the relation and carries it out on the page; what no one has done is to compute both forms and compare them, which is the only part of the claim arithmetic can settle. Measured here over three backgrounds, nine modes $0.3 \le k \le 600$ and $2 \le \ell \le 600$, the two agree to 9.5×10^{-6} worst and 2.7×10^{-8} median against a gate of 10^{-3}. The comparison is stated at the transfer level at fixed k with the same line-of-sight weight on both sides: at C_ℓ the two branches differ by a factor of eighty in the plateau, and that difference is the weight the line-of-sight integral carries — a property of what (102) prints, recorded as an approximation and measured separately, not a disagreement to be gated away.

The external comparison is still run, still stamped, and reported alongside as context. Its value is not a pass mark but an itemisation: how far a $1 + 3$ slow-decoupling solution stands from a Boltzmann code, in bands, with the dominant term named.

The measurement itself is recorded here rather than only in the code. The external codes are not carried as dependencies of the bundle and are not run by the harness: `scripts/make_reference_spectra.py` regenerates the stamped spectra from CAMB and CLASS, `scripts/compare_criterion_5b.py` reproduces the comparison from those stamps with neither code installed, and the table below is what that script writes. It is a committed artefact of the bundle, recoverable and reproducible, and deliberately outside the bundled testing.

Shape, D_ℓ renormalised at $\ell = 30$. Context, not the gate (S4c).

band	vs CAMB	vs CLASS	CAMB vs CLASS
plateau and late ISW, $2 \le \ell < 30$	97.8%	97.8%	—
acoustic peaks, $30 \le \ell < 1000$	137.0%	137.3%	0.370%
First acoustic peak			
this bundle	CAMB	CLASS	
$\ell = 170$	$\ell = 198$	$\ell = 198$	

Itemised: the part this bundle can attribute. The damping coefficient, 0.1% in the plateau and late ISW and 17.2% in the acoustic peaks, between (dif-damp-correct) at $f_2 = 3/4$ and (dif-damp) at $f_2 = 1$ — both *Annals II*'s own. The remainder is the slow-decoupling order, the visibility weight of (102) (approximations A-1 and A-2), and the acoustic scale, $z_* = 1114.5$ here against CAMB's 1088.5.

[6] The reduction of (200) carries a factor of the comoving distance that the displayed (201) does not; substituting $z = k\chi$ gives χ^{2-n}, which is χ^1 at $n = 1$. The companion result (204) descends from a $P(k)$ definition differing by one power of k and is correctly free of that factor, so the asymmetry is not a units convention. A misprint, corrected silently and recorded.

[7] That derivation uses the Bessel identity of §8.3, which requires $[(m/2)!]^2$ in its denominator. As printed it is exact where $\Gamma(m/2 + 1) = 1$, so the large-scale result (201) is unaffected; the horizon-scale normalisation (204) and the D_ℓ of (205) are. A misprint, corrected silently and recorded.

What a self-comparison does not establish is stated here rather than left to be inferred, because it is the obvious objection and it is a fair one. Both routes are handed the same background, the same recombination history, the same acoustic solution, the same potentials and the same damping, and a comparison that holds an input fixed cannot check that input. This is not a formality: setting the gate turned up an under-resolved recombination grid that moves the primary source at η_* by 1.2%, shared by both routes and invisible to the comparison between them. It is recorded, raised, and left for the decision it needs rather than changed inside a release.

The external codes of criterion 5b are compared in `provenance/EXTERNAL-CODE-COMPARISON-v1.md`, which records what CAMB and CLASS each model and — more to the point — what the comparison structurally cannot settle. Neither carries the $1 + 3$ covariant formalism, neither distinguishes mode from multipole coefficients, and neither carries the $O(\varepsilon^2\ell)$ couplings of v1.5.0. So 5b is a check on the *linear* reconstruction and is incapable of validating the nonlinear extension, which is precisely why v1.0.0 must carry no $O(\ell)$ coupling: if one leaked in, the disagreement with CAMB and CLASS would be indistinguishable from a bug. The canonical parameter set switches reionisation, lensing and massive neutrinos **off**, each being physics this bundle does not model.

8 Evidence

Each figure is generated by the script named in its caption and is accompanied by a standalone caption file in `captions/`. This section is itself generated by `scripts/make_figure_pages.py`, so a figure appears here only once it exists in the tree, and what is outstanding is stated rather than omitted.
Deferred, by recorded decision: **F2** truncation convergence, both frames; **F5** source decomposition, two frames; **F7** frame specialisation. Recorded choices, not omissions: the reasoning for each is in `provenance/DECISIONS-v1.0.0.md`. The frame-covariance claim F5 would have drawn is established symbolically in this release, as an exact identity, by `functions/frames.py`. The full specification is `captions/FIGURE-PLAN-v1.0.0.md`.

9 Parameter and numerical evidence

Table 5 lists every number the numerics use. Its *provenance* column distinguishes values taken from the source paper, from standard values that are not, from values derived here, and from values stipulated for this release. The distinction matters: $\Omega_\gamma h^2$ is a standard value and not one the antecedent supplies, and it enters only through $R(\eta_*)$ and hence the sound speed.

Table 5: Every number the numerics use. *Provenance* distinguishes values taken from the source paper, standard values that are not, values derived here, and values stipulated for this release.

Symbol	Code	Value	Provenance	Set in
k_{com}	`K_COM`	1.0	stipulated	`scripts/figure_f1_appendix_f.py`
$\ell_{\max}$	`ELL_MAX`	80	stipulated	`scripts/figure_f1_appendix_f.py`
ℓ_{compared}	`ELL_COMPARED`	25	stipulated	`scripts/figure_f1_appendix_f.py`
rtol, atol	`rtol, atol`	10^{-10}, 10^{-12}	stipulated	`functions/harmonics/free_streaming.py`
z_*	`Z_STAR`	1100	standard	`scripts/derive_ell_range.py`
h	`H`	0.5	AnnalsII §8.3.2	`scripts/derive_ell_range.py`

Symbol	Code	Value	Provenance	Set in
Ω_B	`OMEGA_B`	0.05	AnnalsII §8.3.2	`scripts/derive_ell_range.py`
$\Omega_\gamma h^2$	`OMEGA_GAMMA_H2`	2.47×10^{-5}	standard	`scripts/derive_ell_range.py`
tolerance	`ADOPTED_TOLERANCE`	5%	derived	`scripts/derive_ell_range.py`
acceptance threshold	`--`	10^{-11}	stipulated	`scripts/figure_f1_appendix_f.py`

10 Numerical resolution and reproducibility

The hierarchy is integrated with an explicit Runge–Kutta method of order 8 (`DOP853`) at relative tolerance 10^{-10} and absolute tolerance 10^{-12}, truncated at $\ell_{\max} = 80$ with comparisons restricted to $\ell \leq 25$ so that reflection from the truncation does not contaminate them. The acceptance threshold used in the figures, 10^{-11}, sits an order above the integrator floor.

The release route is

```
python scripts/run_all.py --strict
```

which regenerates every output, figure and table, recomputes the SHA-256 manifests, and reports any artefact that differs from the recorded hash.

Three manifests, and only one of them binds. `PORTABLE-MANIFEST-SHA256.txt` is the release fingerprint: it *excludes the renderings* — everything under `figures/` and `supplementary-materials/`, and every `.png` and `.pdf` — and canonicalises numeric output to nine significant figures, so a machine that did not generate the artefacts can still satisfy it. A mismatch there is release-blocking. `FILE-MANIFEST-SHA256.txt` is then a *within-machine* drift check, reported by artefact class, where only an `exact`-class mismatch blocks. `IMMUTABLE-MANIFEST-SHA256.txt` covers frozen reference material, where any change is release-blocking. The corrections record is append-guarded rather than frozen, because it must take new findings: appending is permitted and rewriting recorded bytes is not.

Why nine PNGs always report as drifted. Every `.png` written into a working copy through the Claude desktop bridge gains a 5 770-byte `caBX` provenance chunk, inserted immediately after `IHDR` by that application and not by anything in this bundle. The pixels are bit-identical — `IHDR`, `pHYs` and every `IDAT` chunk match exactly — and the `.pdf` of the same figure round-trips byte-for-byte, as does every `.csv` and every source file. These nine therefore sit in the `rendered` class, are reported and do not block, which is the rule finding **T-15** wrote for a different cause with the same shape: fonts, freetype and TEX Live versions, where *a PNG that differs in hinting carries no scientific difference.* **The PDF is the authoritative artefact of every figure in this bundle**; the PNG is a convenience copy. Finding **T-36**. The drift is explained here rather than silenced by canonicalising PNGs, because a manifest that stopped reporting a class of change for tidiness would stop reporting real ones too.

Weights are carried as exact rationals rather than floating-point numbers wherever a sign or an ℓ-weight is asserted.

11 Beyond this release: the almost-EGS bounds

The Ehlers–Geren–Sachs theorem states that exact CMB isotropy for all fundamental observers in an expanding dust universe forces exact Friedmann–Lemaître–Robertson–Walker geometry. The *almost*-EGS theorem of Stoeger, Maartens & Ellis, *Ap. J.* **443**, 1 (1995) relaxes that to bounds: limits on the CMB multipoles τ_{A_ℓ} become limits on σ_{ab}, ω_{ab}, $\dot{u}_a$, E_{ab} and H_{ab}.

Those bounds are stated in the *covariant multipoles* — not in mode coefficients — which is what this bundle computes and what a gauge-fixed code cannot supply without a translation. Applying them is

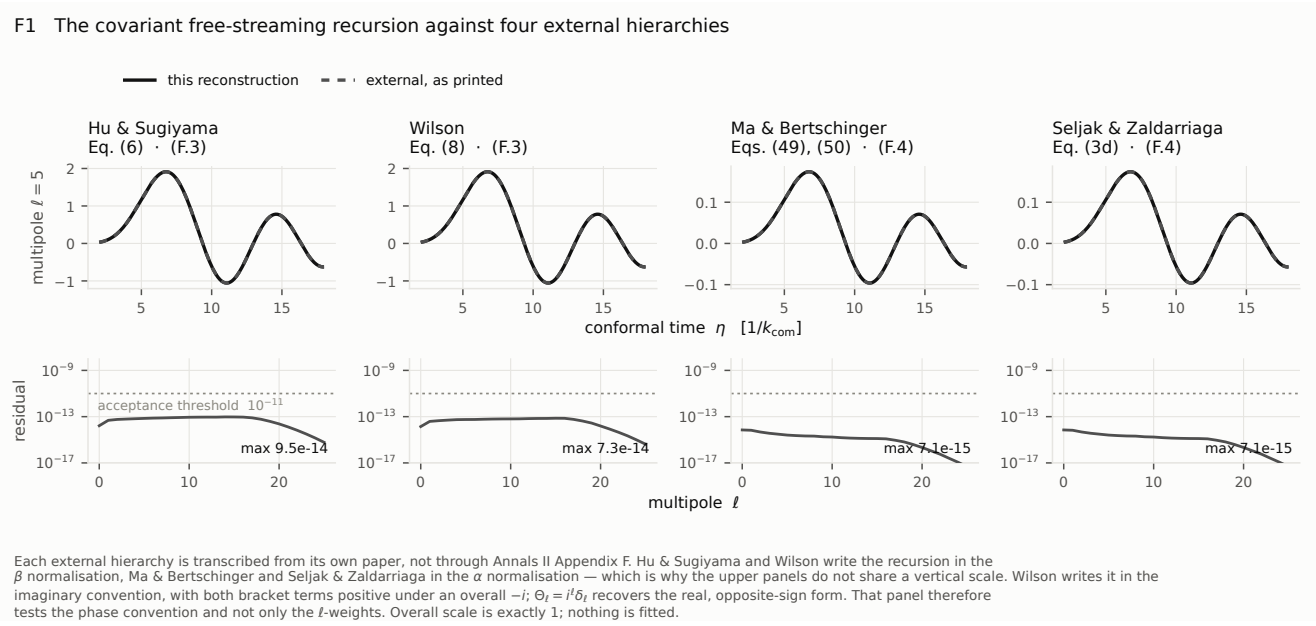


Figure 1: **F1 — The covariant free-streaming recursion against four external hierarchies.** The covariant free-streaming recursion of *Annals II* Appendix F, integrated from a unit monopole, against four external treatments of the same hierarchy. Each is transcribed from its own paper, never through Appendix F. **Upper panels, context:** the reconstruction and the external form at $\ell = 5$; they coincide. The panels do *not* share a vertical scale, because two normalisations are on display. **Lower panels, the claim:** $\max_\eta |\text{external} - \text{reconstruction}|$ against ℓ, with the acceptance threshold marked. Hu & Sugiyama write the recursion in the β normalisation and Ma & Bertschinger and Seljak & Zaldarriaga in the α normalisation, which is why Appendix F prints both (F.3) and (F.4); a basis phase leaking into the couplings would break at least one match. Overall scale is exactly 1.000000000 — nothing is fitted. Four panels from five sources: Wilson & Silk (1981) is the earliest, but at $K = 0$ it is Wilson (1983) Eq. (8) exactly, so it is credited rather than drawn.

Evidence category: cross-check. *Generated by* `scripts/figure_f1_appendix_f.py`. *Standalone caption:* `captions/F1-appendix-f.md`.

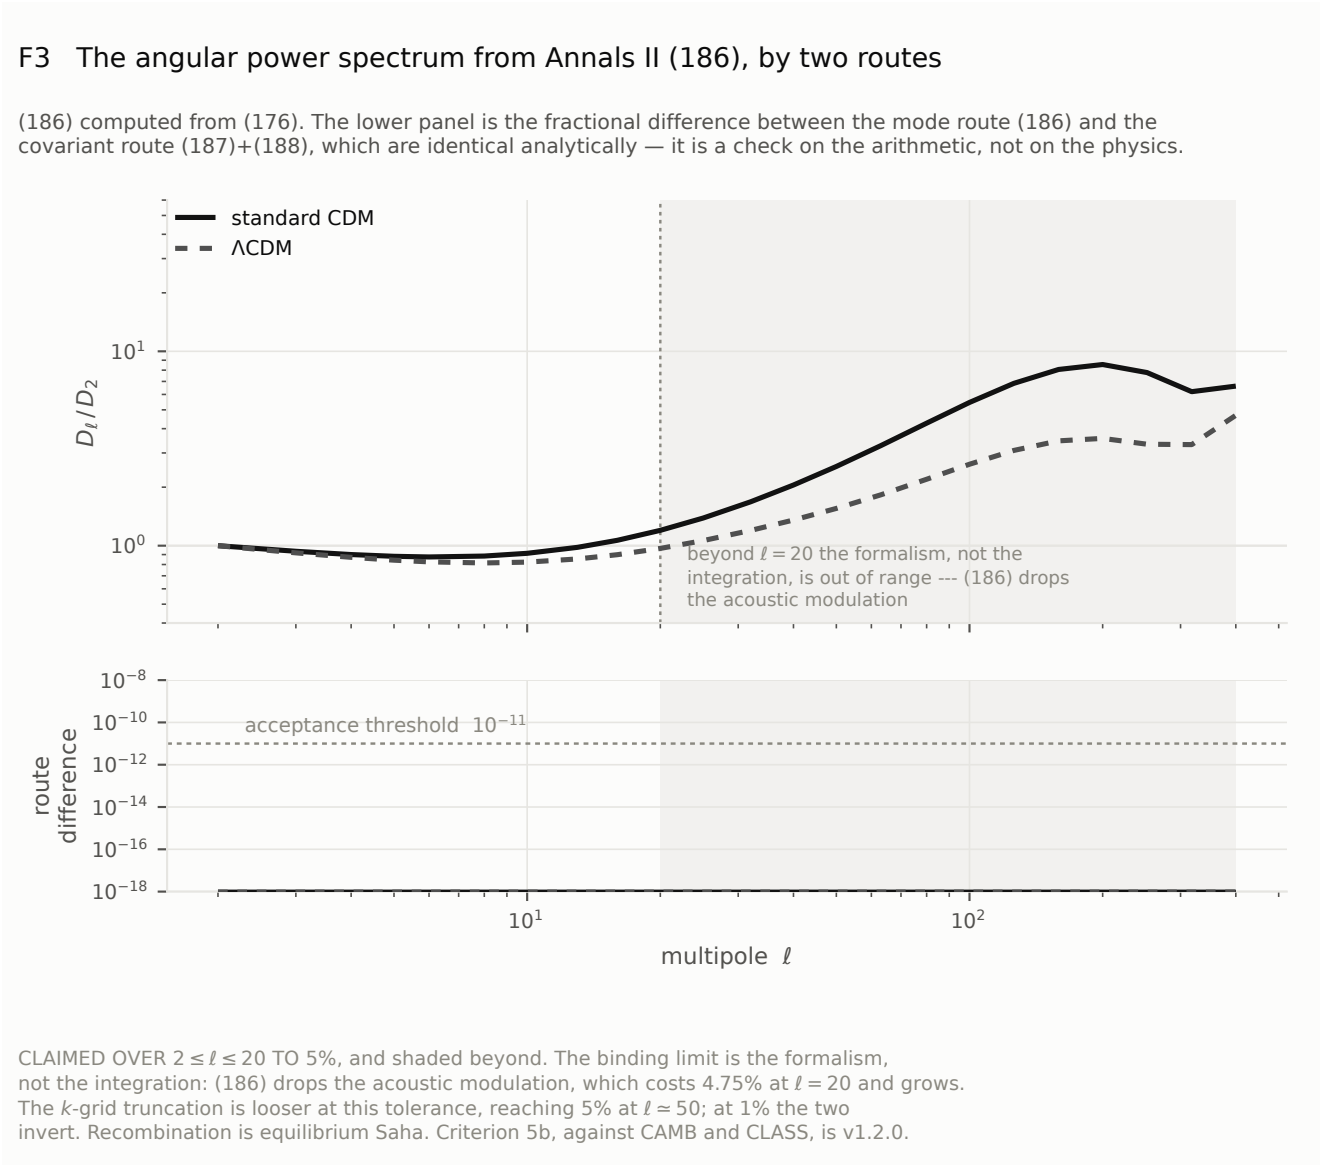


Figure 2: **F3 — The angular power spectrum from Annals II (186), by two routes.** Eq. (186) computed from Eq. (176), for the standard-CDM model of §8.3.2 and for ΛCDM, which *Annals II* does not treat. **Upper panel:** $D_\ell = \ell(\ell+1)C_\ell/2\pi$, each normalised to its own $\ell = 2$ value, so that the comparison is of *shape* and not of an amplitude neither model fixes here. **Lower panel:** the fractional difference between the mode route of (186) and the covariant route of (187) with (188). Those two are identical analytically, so the panel is a check on the arithmetic and not on the physics — and it is worth having, because it caught a factor of four in this bundle's own harmonic weights before any figure was drawn. It sits at 10^{-16}, five orders below the 10^{-11} acceptance threshold.

Evidence category: recovery. *Generated by* `scripts/figure_f3_angular_spectrum.py`. *Standalone caption:* `captions/F3-angular-spectrum.md`.

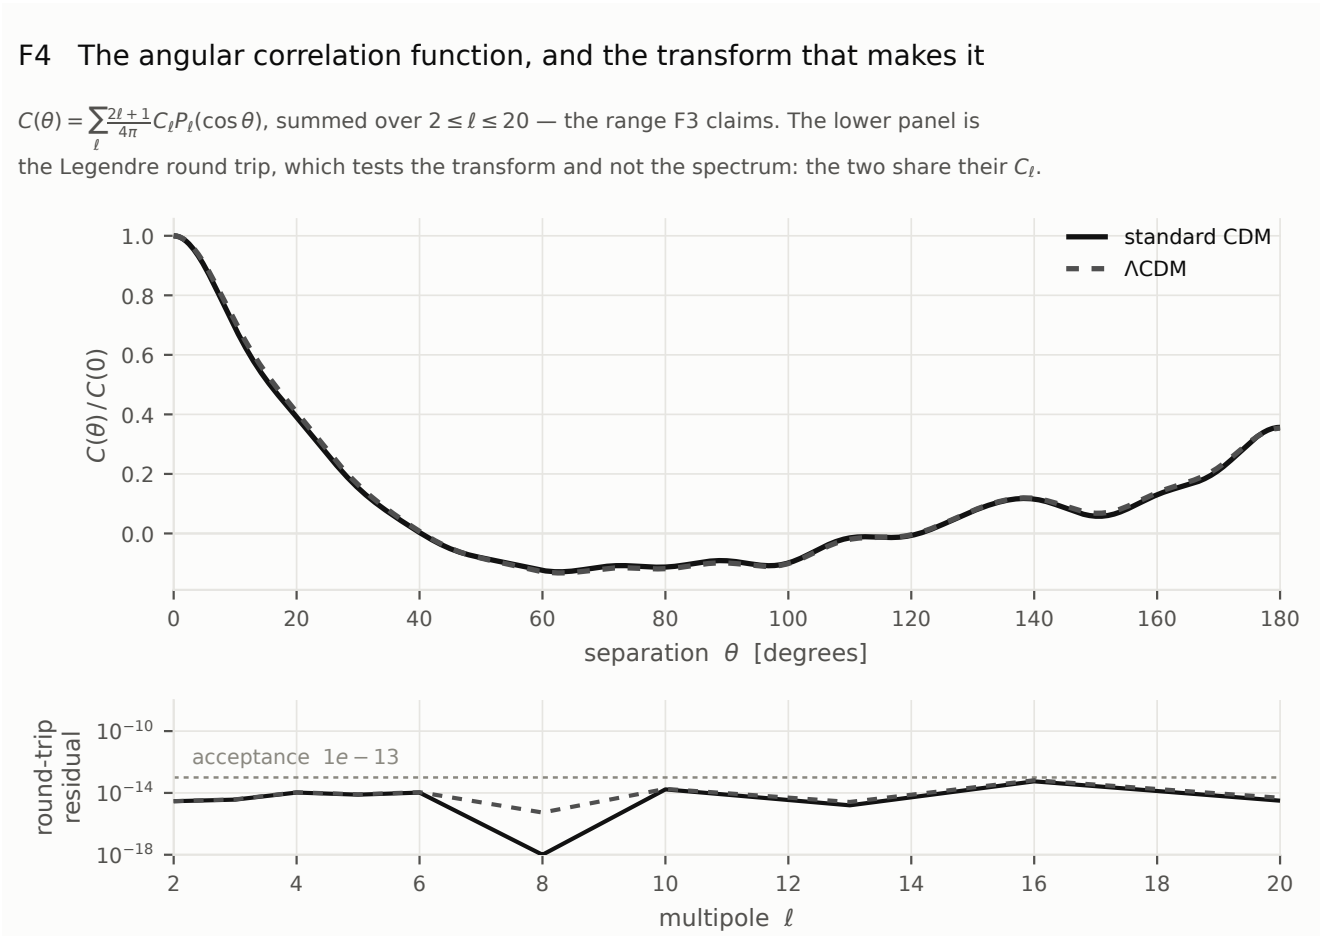


Figure 3: **F4 — The angular correlation function, and the transform that makes it.** $C(\theta) = \sum_\ell \frac{2\ell+1}{4\pi} C_\ell P_\ell(\cos\theta)$, the real-space view of Figure 3: the same content read as a correlation between two directions separated by θ rather than as power at a multipole. **Upper panel:** $C(\theta)/C(0)$ for standard CDM and ΛCDM, summed over $2 \le \ell \le 20$ — the range Figure 3 claims, bounded by (186) dropping the acoustic modulation. Summing further would draw multipoles this release does not stand behind, and a real-space curve hides which ℓ contributed, so the restriction is made in the transform rather than in this caption. **Lower panel:** the Legendre round trip $C_\ell \to C(\theta) \to C_\ell$, at 4×10^{-14} and 7×10^{-14} against an acceptance of 10^{-13}. It tests the transform and *not* the spectrum — the two share their C_ℓ — which is the same caution Figure 3's residual panel carries and for the same reason.

Evidence category: acceptance. *Generated by* `scripts/figure_f4_correlation.py`. *Standalone caption:* `captions/F4-correlation.md`.

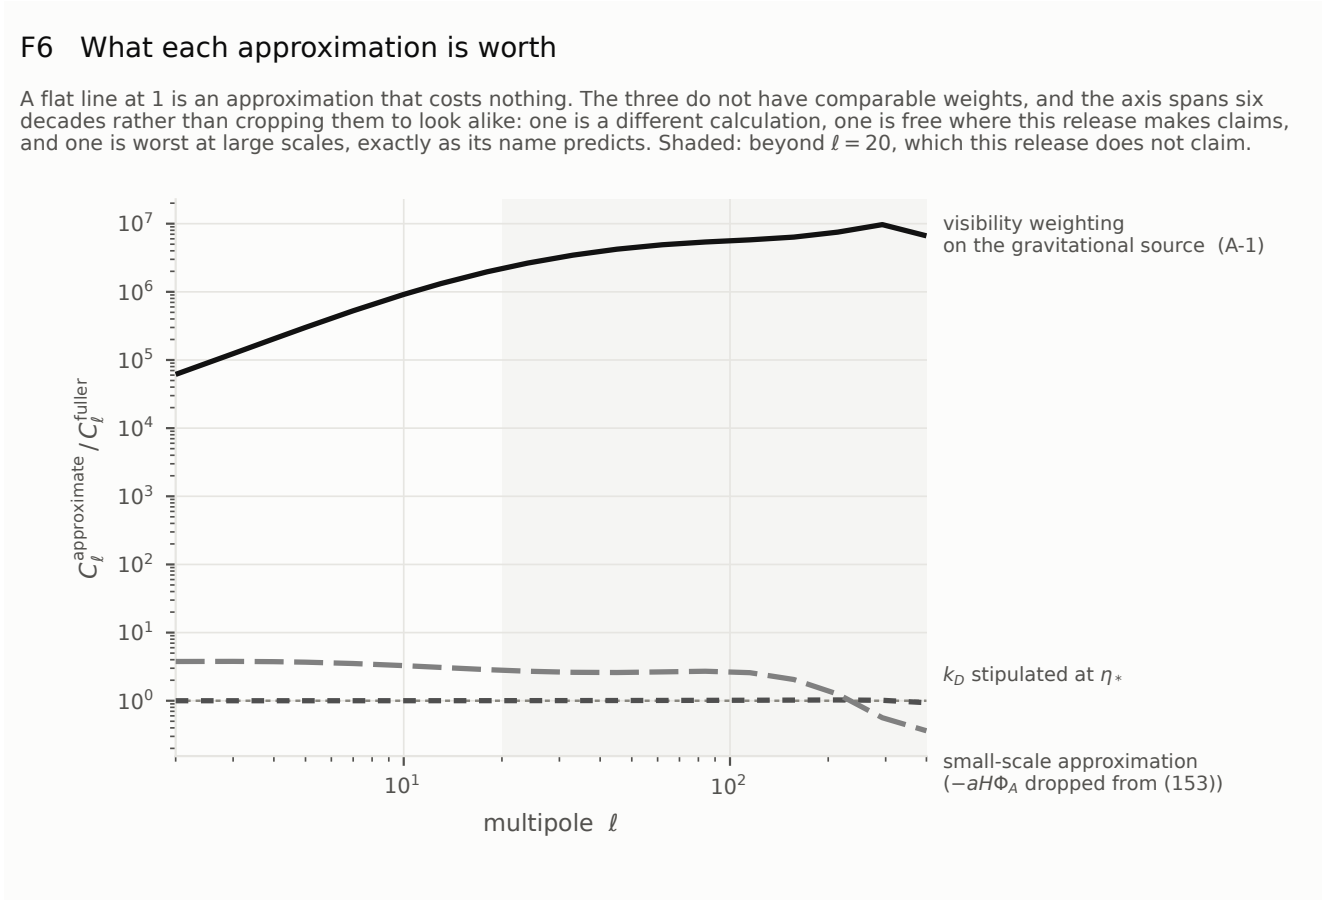


Figure 4: **F6 — What each approximation is worth.** Each series is $C_\ell^{\text{approximate}}/C_\ell^{\text{fuller}}$ for standard CDM, so the panel is scale-free and a flat line at unity is an approximation that costs nothing. *Annals II* makes three narrowing choices which this bundle reproduces rather than corrects, because reproducing a paper means reproducing its approximations (decision S15); they are recorded as **A-n** in `provenance/SOURCE-APPROXIMATIONS-v1.md`, deliberately apart from the defects recorded as **B-n**. This is the one place their cost is a number rather than an adjective.

The three do not have comparable weights, and that is the result, so the axis spans six decades rather than cropping them to look alike. The visibility weighting on the gravitational source (A-1) runs to 10^4–10^6: it is a *different calculation*, not a perturbation of one, which is why `Weighting` has no default. Stipulating k_D at last scattering costs under 0.1% across the claimed range and about 5% by $\ell = 400$ — it **vindicates the release**, which takes the cheap choice where it makes claims. The small-scale approximation, dropping $-aH\Phi_A$ from (153) and so producing the printed oscillator (154), departs from unity most at *large* scales, which is what its name predicts.

A correction kept because it nearly became this figure. The k_D series was first computed with $k_D(\eta)$ along the whole ray. After last scattering $\kappa' \to 0$, so $1/\kappa'$ in (dif-damp) diverges and k_D collapses, damping every integrated source to nothing — a plasma formula evaluated in a universe that no longer has a plasma. That version reported the release's choice as costing a factor of nine at $\ell = 2$; the factor was the artefact. Clamped to $k_D(\min(\eta, \eta_*))$ the comparison is a real one. Shaded: beyond $\ell = 20$, which this release does not claim.

Evidence category: diagnostic. *Generated by* `scripts/figure_f6_approximations.py`. *Standalone caption:* `captions/F6-approximations.md`.

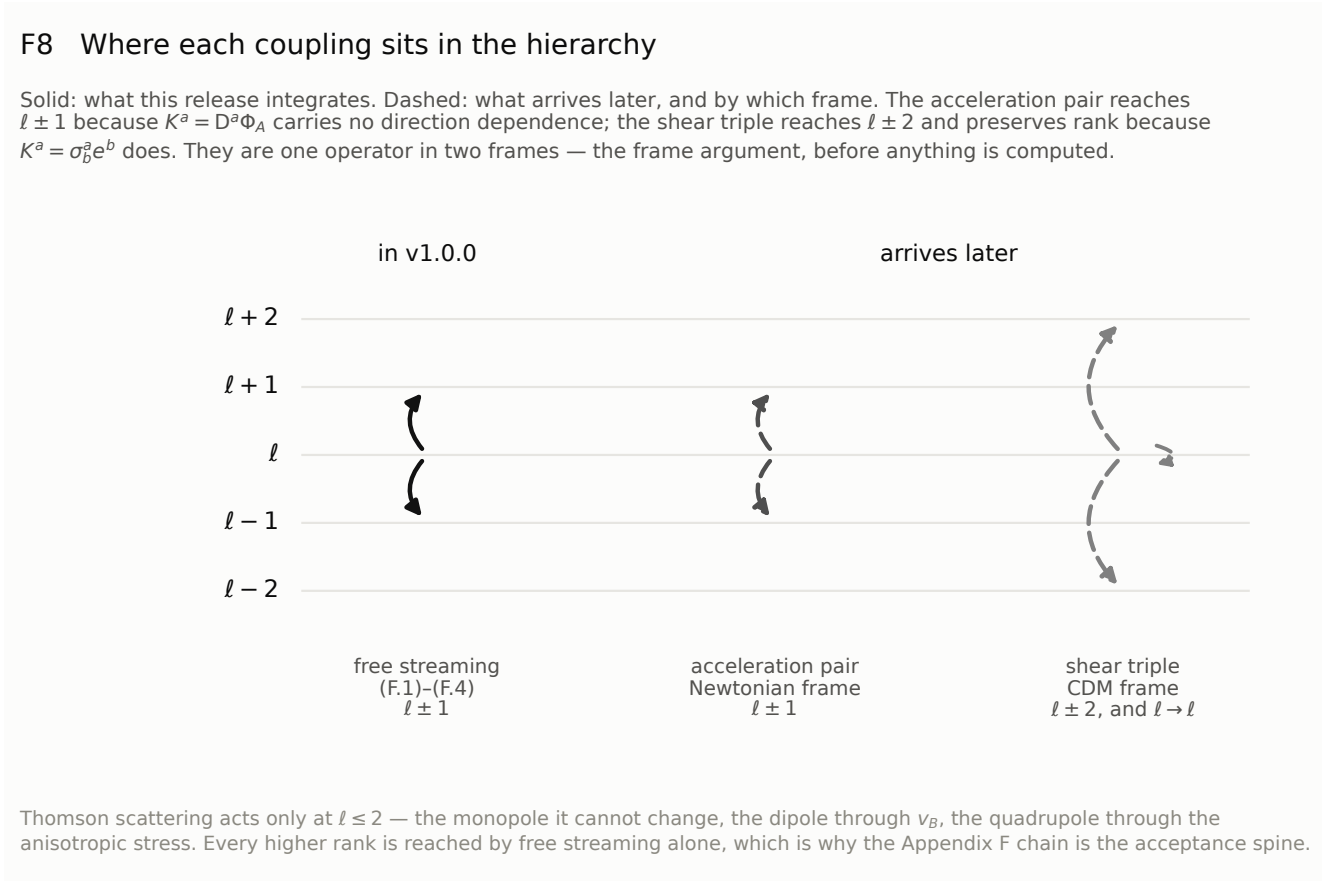


Figure 5: **F8 — Where each coupling sits in the hierarchy.** The one released figure that computes nothing: a statement about the *structure* of the exact divergence hierarchy, read off the equations rather than measured. It is generated by a script like every other figure, so that the manifest covers it — a hand-drawn diagram is the one artefact that can drift without any gate noticing.

Solid: what this release integrates. Free streaming, (F.1)–(F.4), couples rank ℓ to $\ell \pm 1$ and to nothing else, which is the whole of v1.0.0's hierarchy and the reason the Appendix F chain is the acceptance spine. **Dashed:** what arrives later, and by which frame. The acceleration pair reaches $\ell \pm 1$ because $K^a = \mathrm{D}^a\Phi_A$ carries no direction dependence of its own; the shear triple reaches $\ell \pm 2$ *and* preserves rank because $K^a = \sigma^a{}_b e^b$ does. The rank-preserving coupling exists only when K^a depends on direction, and is the single structural difference between the two frames.

Both are the same operator $K^a_\perp \partial/\partial e^a$, and the number of ranks each reaches is fixed by that one property. **That is the whole content of the frame argument, visible before any of it is computed.** Decision S9 is drawn as well as asserted: every coupling in the solid half is a ratio of linear polynomials and tends to a constant, so an $O(\ell)$ coupling wired in early would fail `tests/test_appendix_f_chain.py` at large ℓ.

Evidence category: schematic. *Generated by* `scripts/figure_f8_schematic.py`. *Standalone caption:* `captions/F8-schematic.md`.

therefore natural work for this formalism, and it is **v2.0.0** work, because it is a calibration against data rather than an analytical reconstruction.

Table 6: The structure v2.0.0 fills. **No cell is computed at v1.0.0**, and none is quoted here from memory: observational columns are filled from the mission papers at v2.0.0 under the same rule that equation numbers are resolved from the source.

Mission	ℓ **reach**	**bound, standard CDM**	**bound, ΛCDM**
COBE-DMR (1992)	$\lesssim 20$	*the original*	*not in the original*
WMAP (2003–2013)	$\lesssim 1000$	pending	pending
Planck (2013–2018)	$\lesssim 2500$	pending	pending

Why the table needs both columns. The original bounds were derived for a CDM model, not for ΛCDM, and the difference is not cosmetic. The argument converts multipole limits into limits on σ/H using the background expansion, and Λ changes $H(z)$ at low redshift. More sharply, Λ adds a source the original did not have: the late integrated Sachs–Wolfe effect contributes to exactly the low-ℓ multipoles the COBE bounds rest on. In standard CDM it is absent — Φ is constant, which this bundle's growth factor reproduces to six figures — and in ΛCDM it is not, Φ falling to 0.779 of its early value by today. A low-ℓ limit therefore constrains a *different combination* in the two models, so a bound quoted for one is not a bound for the other. The CDM column has a published precedent; the ΛCDM column is new.

Beyond v2.0.0. The published bounds are first order in the anisotropy parameter. The natural extension is the next order, and this project is placed to make it because **v1.5.0 supplies the $O(\varepsilon^2\ell)$ couplings a second-order argument needs** and no external code carries them. Stated as a direction and not a plan: it is gated on v1.5.0, and the correction's sign and size are unknown until it is done. It is recorded so the v2.0.0 table is built in a shape that can take it. See `provenance/ALMOST-EGS-BOUNDS-v1.md`.

12 Interpretive limits

This is an analytical reproducibility bundle for the associated paper and its supplement. It is not an empirical calibration, not a parameter-estimation pipeline, and not a Boltzmann code for general use. No new effect is claimed and none is reported.

Annals II's own §8.3.2 model is standard CDM — $h = 0.5$, $\Omega_B = 0.05$, flat and matter dominated — and the C_ℓ it prints there is a closed-form Sachs–Wolfe result on large scales with the transfer function set to unity. This release goes past that in one respect and one only: it computes the spectrum from the hierarchy rather than from the closed form, and it computes it for the ΛCDM model as well, which *Annals II* does not treat. The marginal cost is the expansion law and the late-time integrated term; every other ingredient — recombination, tight coupling, diffusion damping — is already required by the CDM case and is not additional. No new *physics* is introduced by Λ, and none is claimed.

What the release does **not** do is calibrate. Comparison is made against CMBFAST and against nothing else, over a derived range, as a check that the reconstruction is right — not as a fit. Parameter estimation and calibration against data are the **v2.0.0** deliverable.

Two statements must travel together, and neither stands alone: the nonlinear contribution $\delta\tau_{NL}$ **matters** — it is not negligible and not an artefact — and it **adds nothing**, redistributing power rather than creating it, because $\delta\tau_{NL} \propto \Phi\Phi$ makes its contribution to C_ℓ quartic in Φ while the linear C_ℓ is quadratic.